



**LET'S
BUILD**
JANUARY 4TH 2020

Machine Learning Best Practices

Shweta Bhatt | ML GDE



Agenda

- Motivation
- ML Workflow
 - Data Collection and Preparation
 - Model Selection and Training
 - Model Validation and Error Analysis
 - Testing and Deployment
 - Retraining and Maintaining Models
- Bias vs Variance trade-off
- General best practices
- Summary

Why do we need ML best practices?

- Labelled vs unlabelled data?
- Feature Selection?
- Model Architecture?
- Model Complexity?
- Metrics?
- Hyperparameter Tuning
- Next Iteration?
- Retraining?



[Image source](#)

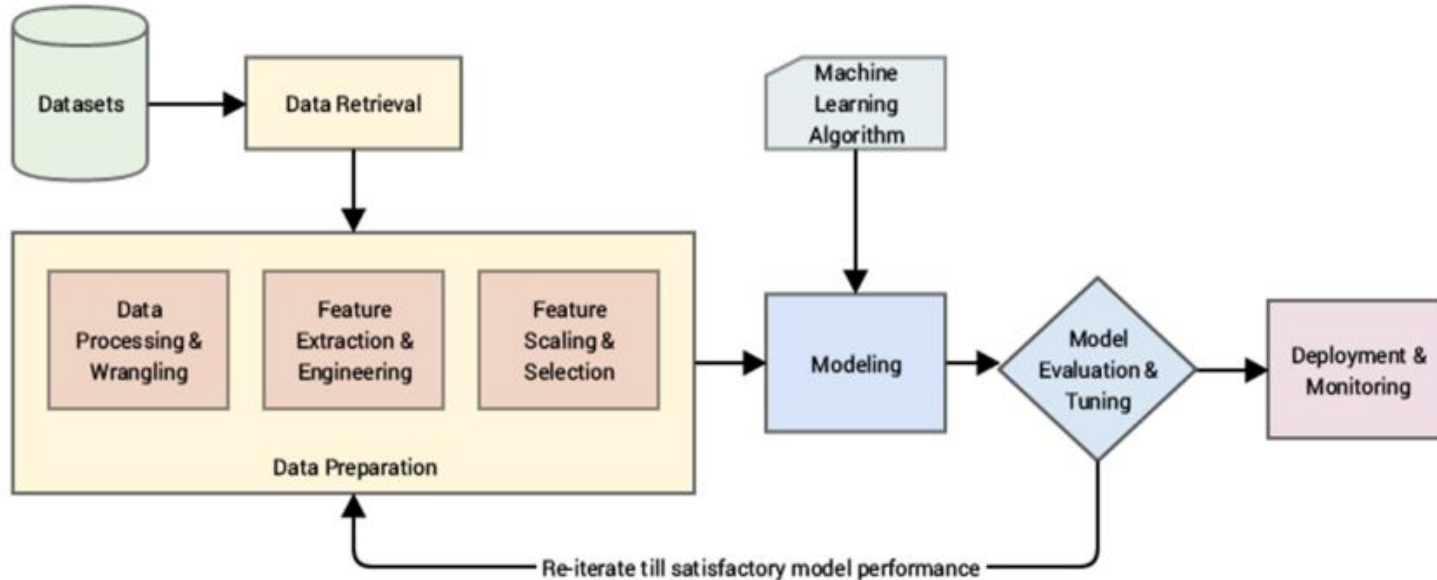
Our goal

Discuss some best practices that would aid in providing a direction that we can follow while making decisions when we are building ML systems.

ML Workflow: summary

- Problem Understanding and Formulation


-



- Maintenance and Active Learning

[Image source](#)

Why is problem formulation important?



[Image source](#)

Problem Understanding and Formulation

- Define what ML problem to solve given the objectives, resources (data) and constraints
- Inputs and outcome of the model
- Metrics: ML, Business KPI
- Heuristics vs ML
- Formulation: start with a simple baseline, keep adding complexity later if required

It all starts with data



Image source

Data Collection and Preparation

- Collection: Scraping vs using an API
- Storage: Versioning, Database platform
- Processing: Cleaning + Formatting
- Verification and Validness: Consistency + Completeness

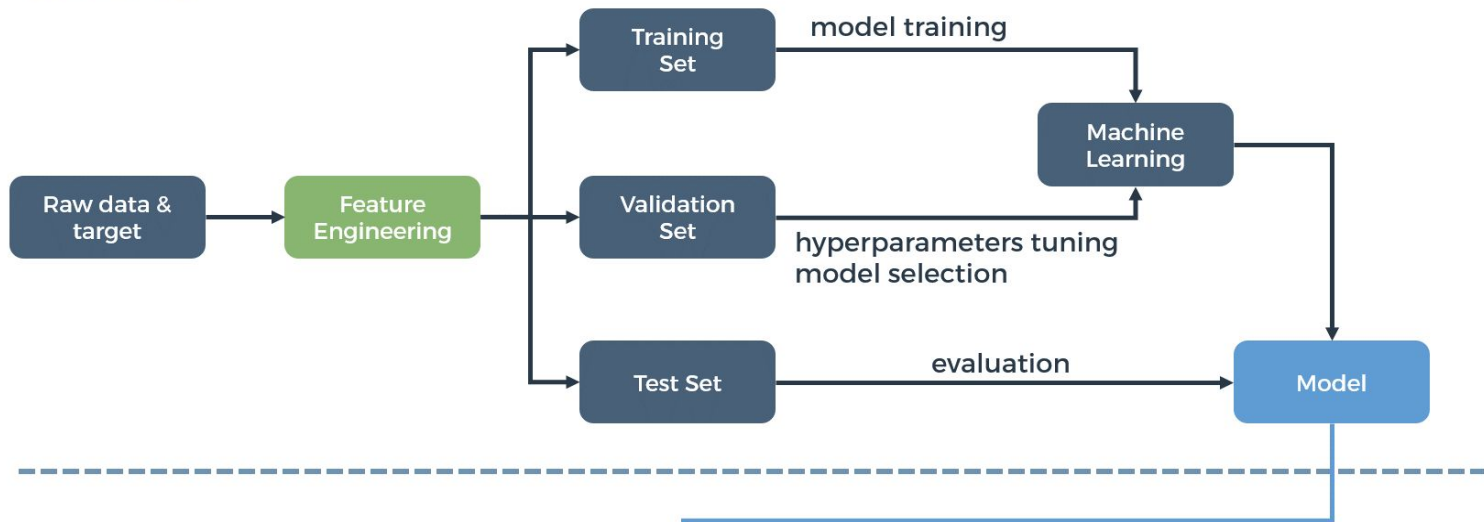
Where would you find an item quicker?



[Image source](#)

Training and prediction pipelines

TRAINING



PREDICTING



Organizing Codebase

- Directory Structure
- Saving model names with parameters
- Experimentation notebooks
- Data Preprocessing
- Training scripts
- Inference/Prediction Scripts


| | |
|--------------------|---|
| — LICENSE | |
| — Makefile | <- Makefile with commands like `make data` or `make train` |
| — README.md | <- The top-level README for developers using this project. |
| — data | |
| — external | <- Data from third party sources. |
| — interim | <- Intermediate data that has been transformed. |
| — processed | <- The final, canonical data sets for modeling. |
| — raw | <- The original, immutable data dump. |
| — docs | <- A default Sphinx project; see sphinx-doc.org for details |
| — models | <- Trained and serialized models, model predictions, or model summaries |
| — notebooks | <- Jupyter notebooks. Naming convention is a number (for ordering), the creator's initials, and a short `-` delimited description, e.g. `1.0-jqp-initial-data-exploration`. |
| — references | <- Data dictionaries, manuals, and all other explanatory materials. |
| — reports | <- Generated analysis as HTML, PDF, LaTeX, etc. |
| — figures | <- Generated graphics and figures to be used in reporting |
| — requirements.txt | <- The requirements file for reproducing the analysis environment, e.g. generated with `pip freeze > requirements.txt` |

- └─ setup.py <- Make this project pip installable with `pip install -e`
- └─ src <- Source code for use in this project.
 - └─ __init__.py <- Makes src a Python module
 - └─ data <- Scripts to download or generate data
 - └─ make_dataset.py
 - └─ features <- Scripts to turn raw data into features for modeling
 - └─ build_features.py
 - └─ models <- Scripts to train models and then use trained models to make
 - predictions
 - └─ predict_model.py
 - └─ train_model.py
 - └─ visualization <- Scripts to create exploratory and results oriented visualizations
 - └─ visualize.py
- └─ tox.ini <- tox file with settings for running tox; see tox.testrun.org

[source](#)

Building an image classifier (apple vs orange)

- Large labelled dataset
- Train-test split: 70%-30%
- Testing accuracy: 95%
- Accuracy on the app (real world data) : 70%
- **What went wrong?**



[Image source](#)

How to setup dev and test sets?

- **Distribution**
 - Should reflect the data you expect to get in future (unseen data)
 - Should come from the same distribution
- **Size**
 - Should be large enough to detect differences between different algorithms
 - 30% of the dataset vs 1000-10000 examples
- **Effort**
 - Come up with dev and test sets quickly (if not mature applications)
 - Change them quickly if you realize they are not meeting the mark

Model Building

- Feature Engineering & Selection
- Model Selection - follow general rule of thumb
- Model Training - early stopping, schedule training process accordingly, debugging and investigating, for example: use Tensorboard if you are using Tensorflow for implementation
- Model Validation - measure performance against a benchmark
- Visualizations - examine learning curves, visualize metrics

What other approaches to try?




[Image source](#)

Error Analysis

- Take a sample of 100 examples from your dev set that were misclassified by your model
- Evaluate these manually to understand the underlying causes of these errors
- Make a list of any patterns among the errors and try to categorize them
- Advantages:
 - You can evaluate how promising different directions and prioritize your ideas accordingly
 - Can save you a lot time and effort

How do you select the best model?




[Image source](#)

Evaluation Metrics

- Appropriate metrics for context and objectives of the system
- **Single-number evaluation metric** \Rightarrow decide which model works best
- **Derived metric**
 - E.g. F1 score, weighted average
- If N multiple metrics cannot be combined directly
 - **N-1 satisfying metrics** \Rightarrow they meet a certain value
 - **1 optimizing metric** \Rightarrow maximize performance over this
 - E.g. running time - satisfying metric, accuracy - optimizing metric

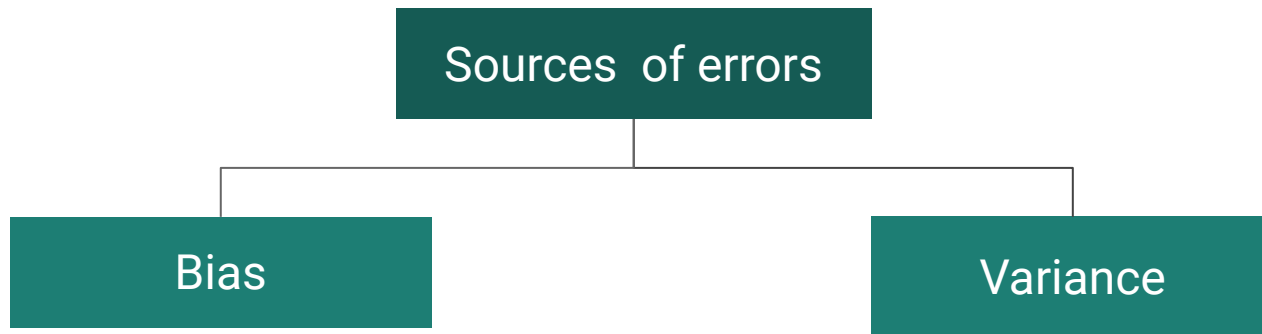
Iterative Process

- Don't try to build the perfect solution in one go
- Start with something simple as quickly as possible
- Use error analysis to suggest you promising directions
- Use that to improve your existing solution




[Image source](#)

Bias vs Variance tradeoff



Our goal is to optimize for **low bias** and **low variance**


What do you interpret from this curve?



How to reduce bias?

- Increase model complexity (no. of layers/neurons) if computational power is not a limitation
- Update input features based on the feedback received from error analysis
- Reduce regularization
- Change the model architecture


Again, what do you interpret from this curve?



How to reduce variance?

- Add more data
- Add regularization
- Early stopping
- Decrease number of input features
- Decrease model complexity (number of layers/neurons)
- Update input features based on the feedback received from error analysis
- Change the model architecture

And this one?



Testing and Deployment

- Unit testing - qualitative and quantitative
- Virtual Environment Setup
- Containers - E.g. Docker
- Do not reinvent the wheel - use open source code + tools


Maintenance and Active Learning

- Performance checks
- When to retrain?
- How to incorporate new data?

Human-centered Design Approach

- Clarity and control
- Engage with diverse set of users
- Incorporate feedback before and through project deployment

Examine raw data \Rightarrow Exploratory Data Analysis



[Image source](#)

Understanding limitations

- What does your model solve?
- Where does it fall short?
- Can we overcome that with some heuristics?
- Communication with the team member and stakeholders

Test, test, test

- Unit tests
- Integration tests
- Update gold standard datasets
- Quality checks

System monitoring and update

- Real world performance + feedback
- Short term vs long term solution
- When to update?
- Effects of updating: system quality, user experience



[Image source](#)

Summary

- Understand the problem well and formulate it correctly
- Define dev and test sets and metrics
- Start with a simple baseline and add complexity when required
- Understand the data and clean it adequately
- Use feedback from data and results \Rightarrow Error analysis
- Do not reinvent the wheel \Rightarrow Use open source libraries, code, tools, literature
- Carry out a quick exhaustive research before jumping to implementation
- Use visualizations to make your life better
- Test your systems thoroughly
- Deploy as per your need \Rightarrow web/mobile

Thank you

Questions?

Let's connect-



[@shweta_bhatt8](https://twitter.com/@shweta_bhatt8)



[@shweta_bhatt](https://medium.com/@shweta_bhatt)

References

- [Machine Learning Yearning by Andrew Ng](#)
- [Best Practices in Machine Learning Infrastructure-Algorithmia](#)